C1_W2_Lab05_Sklearn_GD_Soln

October 23, 2024

1 Optional Lab: Linear Regression using Scikit-Learn

There is an open-source, commercially usable machine learning toolkit called scikit-learn. This toolkit contains implementations of many of the algorithms that you will work with in this course.

1.1 Goals

In this lab you will: - Utilize scikit-learn to implement linear regression using Gradient Descent

1.2 Tools

You will utilize functions from scikit-learn as well as matplotlib and NumPy.

```
[1]: import numpy as np
     import matplotlib.pyplot as plt
     from sklearn.linear_model import SGDRegressor
     from sklearn.preprocessing import StandardScaler
     from lab_utils_multi import  load_house_data
     from lab_utils_common import dlc
     np.set_printoptions(precision=2)
     plt.style.use('./deeplearning.mplstyle')
```

2 Gradient Descent

Scikit-learn has a gradient descent regression model sklearn.linear_model.SGDRegressor. Like your previous implementation of gradient descent, this model performs best with normalized inputs. sklearn.preprocessing.StandardScaler will perform z-score normalization as in a previous lab. Here it is referred to as 'standard score'.

2.0.1 Load the data set

```
[2]: X_train, y_train = load_house_data()
     X_features = ['size(sqft)','bedrooms','floors','age']
```

2.0.2 Scale/normalize the training data

```
[3]: scaler = StandardScaler()
     X_norm = scaler.fit_transform(X_train)
     print(f"Peak to Peak range by column in Raw        X:{np.ptp(X_train,axis=0)}")
     print(f"Peak to Peak range by column in Normalized X:{np.ptp(X_norm,axis=0)}")
```

```
Peak to Peak range by column in Raw        X:[2.41e+03 4.00e+00 1.00e+00
9.50e+01]
Peak to Peak range by column in Normalized X:[5.85 6.14 2.06 3.69]
```

2.0.3 Create and fit the regression model

```
[4]: sgdr = SGDRegressor(max_iter=1000)
     sgdr.fit(X_norm, y_train)
     print(sgdr)
     print(f"number of iterations completed: {sgdr.n_iter_}, number of weight␣
      ↪updates: {sgdr.t_}")
```

```
SGDRegressor(alpha=0.0001, average=False, early_stopping=False, epsilon=0.1,
             eta0=0.01, fit_intercept=True, l1_ratio=0.15,
             learning_rate='invscaling', loss='squared_loss', max_iter=1000,
             n_iter_no_change=5, penalty='l2', power_t=0.25, random_state=None,
             shuffle=True, tol=0.001, validation_fraction=0.1, verbose=0,
             warm_start=False)
number of iterations completed: 120, number of weight updates: 11881.0
```

2.0.4 View parameters

Note, the parameters are associated with the *normalized* input data. The fit parameters are very close to those found in the previous lab with this data.

```
[5]: b_norm = sgdr.intercept_
     w_norm = sgdr.coef_
     print(f"model parameters:                   w: {w_norm}, b:{b_norm}")
     print( "model parameters from previous lab: w: [110.56 -21.27 -32.71 -37.97], b:
      ↪ 363.16")
```

```
model parameters:                   w: [110.1  -21.04 -32.44 -38.06], b:[363.16]
model parameters from previous lab: w: [110.56 -21.27 -32.71 -37.97], b: 363.16
```

2.0.5 Make predictions

Predict the targets of the training data. Use both the `predict` routine and compute using w and b.

```
[6]: # make a prediction using sgdr.predict()
     y_pred_sgd = sgdr.predict(X_norm)
     # make a prediction using w,b.
     y_pred = np.dot(X_norm, w_norm) + b_norm
     print(f"prediction using np.dot() and sgdr.predict match: {(y_pred ==␣
      ↪y_pred_sgd).all()}")

     print(f"Prediction on training set:\n{y_pred[:4]}" )
     print(f"Target values \n{y_train[:4]}")
```

```
prediction using np.dot() and sgdr.predict match: True
Prediction on training set:
[295.18 485.9  389.61 492.06]
Target values
[300.   509.8 394.   540. ]
```

2.0.6 Plot Results

Let's plot the predictions versus the target values.

```
[7]: # plot predictions and targets vs original features
     fig,ax=plt.subplots(1,4,figsize=(12,3),sharey=True)
     for i in range(len(ax)):
         ax[i].scatter(X_train[:,i],y_train, label = 'target')
         ax[i].set_xlabel(X_features[i])
         ax[i].scatter(X_train[:,i],y_pred,color=dlc["dlorange"], label = 'predict')
     ax[0].set_ylabel("Price"); ax[0].legend();
     fig.suptitle("target versus prediction using z-score normalized model")
     plt.show()
```


3

2.1 Congratulations!

In this lab you: - utilized an open-source machine learning toolkit, scikit-learn - implemented linear regression using gradient descent and feature normalization from that toolkit

[ ]: